

Lecture 7: Optimizing CUDA Programs Part 2: Convolution & MatMul


**Programming Massively Parallel Multiprocessors and
Heterogeneous Systems (Understanding and programming the
devices powering AI)**

Jonathan Appavoo

Convolution & CUDA

on-SM memory
"Use the ~~force~~ Luke"

- Exploit the on-SM memories to magnify arithmetic intensity (mitigate GMEM Latency)
 - Constant Memory can make a big difference
 - What about the L1 and L2
 - truth is, you have to be smart
 - There are plenty of other approaches (eg, register the kernel, embed in instructions, exploit L1, ...
 - not clear who knows (or can know) what is "optimal"
- Until you measure, nothing is real!




Matrix Multiplication

WARNING: "Any sufficiently advanced technology is indistinguishable from magic.", Arthur C Clarke, 1969

Matrix Multiplication: highly parallelizable!

- $C = A \times B$: key computation within many scientific applications particularly those in deep learning
- Compute: $O(N^3)$ -- Memory $O(N^2)$
 - high compute intensity
- each element C dot product of a row of A and a col of B

$$c_{i,j} = \sum_{k=0}^{m-1} a_{i,k} b_{k,j}$$



GPU challenge: minimize GMEM accesses

How many GMEM accesses?


```
for (int i = 0; i < M; ++i)
    for (int j = 0; j < N; ++j)
        for (int k = 0; k < K; ++k)
            C[i][j] += A[i][k] * B[k][j];
```

For square matrices:

$$3 * N^3$$

Idea: use SMEM


?





Idea: use SMEM

- But: A & B & C will not fit in SMEM ☹
- Idea: rearrange loops

```
for (int i = 0; i < M; ++i)
    for (int j = 0; j < N; ++j)
        for (int k = 0; k < K; ++k)
            C[i][j] += A[i][k] * B[k][j];
```






```
for (int k = 0; k < K; ++k)
    for (int i = 0; i < M; ++i)
        for (int j = 0; j < N; ++j)
            C[i][j] += A[i][k] * B[k][j];
```



- copy col of A & row of B into SMEM
- compute its outer product and accumulate in C
- col of A & row of B never used again
- But all of C accessed many time → doesn't fit in SMEM

Idea: use SMEM


- Compute Tiles of C at a time
- Use tiles that fit in SMEM
- ...



Tiled Computation


```
for (int m = 0; m < M; m += Mtile) // iterate over M dim  
    for (int n = 0; n < N; n += Ntile) // iterate over N dim  
        for (int k = 0; k < K; ++k)  
            for (int i = 0; i < Mtile; ++i) // compute one tile  
                for (int j = 0; j < Ntile; ++j) {  
                    int row = m + i;  
                    int col = n + j;  
                    C[row][col] +=  
                        A[row][k] * B[k][col];  
                }
```

- For each C-tile:
tiles of A & B fetched
exactly once
- Ensure the 3 tiles
fit in SMEM




Further considerations

- Decide what each warp does:



Further considerations II

- Optimize warps:



Bottom line for Matrix Multiplication: (and maybe any core routines) Abracadabra!

20.10.4. Features Accelerating Specialized Computations

The NVIDIA Blackwell GPU architecture extends features to accelerate matrix multiply-accumulate (MMA) from the NVIDIA Hopper GPU architecture.

See the [PTX ISA](#) for more details.

This feature set is only available within the CUDA compilation toolchain through inline PTX.

It is strongly recommended that applications utilize this complex feature set through CUDA-X libraries such as cuBLAS, cuDNN, or cuFFT.

It is strongly recommended that device kernels utilize this complex feature set through [CUTLASS](#), a collection of CUDA C++ template abstractions for implementing high-performance matrix-multiplication (GEMM) and related computations at all levels and scales within CUDA.

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#features-accelerating-specialized-computations-12-0>

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#features-accelerating-specialized-computations-10-0>

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#features-accelerating-specialized-computations>

A statement (disclaimer) in the CUDA Programming Guide
that accompanies ALL Compute Capabilities 9 and beyond.

- Surprisingly non-trivial
- Hard to get right from a performance PoV
- Nvidia CUBLAS cannot be beat!
- CUTLESS library is closest to CUBLAS
 - See <https://github.com/NVIDIA/cutlass>