



GPU Teaching Kit  
Accelerated Computing



## Module 10.1 – 10.4 Parallel Computation Patterns (scan)

Prefix Sum, A Work-Inefficient scan kernel, A Work-Efficient Parallel Scan Kernel, & More on Parallel Scan

# Part 1: Prefix Sum

- To master parallel scan (prefix sum) algorithms
  - Frequently used for parallel work assignment and resource allocation
  - A key primitive in many parallel algorithms to convert serial computation into parallel computation
  - A foundational parallel computation pattern
  - Work efficiency in parallel code/algorithms
- Reading –Mark Harris, Parallel Prefix Sum with CUDA
  - [https://developer.nvidia.com/gpugems/GPUGems3/gpugems3\\_ch39.html](https://developer.nvidia.com/gpugems/GPUGems3/gpugems3_ch39.html)

# Inclusive Scan (Prefix-Sum) Definition

**Definition:** *The scan operation takes a binary associative operator  $\oplus$  (pronounced as circle plus), and an array of  $n$  elements*

$$[x_0, x_1, \dots, x_{n-1}],$$

*and returns the array*

$$[x_0, (x_0 \oplus x_1), \dots, (x_0 \oplus x_1 \oplus \dots \oplus x_{n-1})].$$

**Example:** If  $\oplus$  is addition, then scan operation on the array would return

$$\begin{aligned}[3 & 1 & 7 & 0 & 4 & 1 & 6 & 3], \\ [3 & 4 & 11 & 11 & 15 & 16 & 22 & 25].\end{aligned}$$

# An Inclusive Scan Application Example

- Assume that we have a 100-inch sandwich to feed 10 people
- We know how much each person wants in inches
  - [3 5 2 7 28 4 30 8 1]
- How do we cut the sandwich quickly?
- How much will be left?
  
- Method 1: cut the sections sequentially: 3 inches first, 5 inches second, 2 inches third, etc.
  
- Method 2: calculate prefix sum:
  - [3, 8, 10, 17, 45, 49, 52, 52, 60, 61] (39 inches left)

# Typical Applications of Scan

- Scan is a simple and useful parallel building block
  - Convert recurrences from sequential:

```
out[0] = f(0);  
for(j=1; j<n; j++)  
    out[j] = out[j-1] + f(j);
```
  - Into parallel:

```
forall(j) { temp[j] = f(j) };  
scan(out, temp);
```
- Useful for many parallel algorithms:
  - Radix sort
  - Quicksort
  - String comparison
  - Lexical analysis
  - Stream compaction
  - Polynomial evaluation
  - Solving recurrences
  - Tree operations
  - Histograms, ....

# A Work Efficient C Implementation

```
void scanAdd(int *x, int y*, int n) {  
    y[0] = x[0];  
    for (int i = 1; i < n; i++) y[i] = y [i-1] + x[i];  
}
```

Computationally efficient:

N additions needed for N elements - O(N)!

Only slightly more expensive than sequential reduction.

# A Naïve Inclusive Parallel Scan

- Assign one thread to calculate each  $y$  element
- Have every thread to add up all  $x$  elements needed for the  $y$  element

$$y_0 = x_0$$

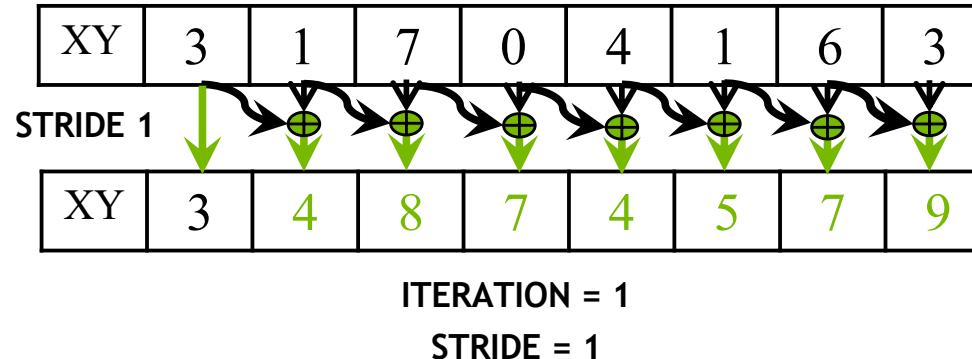
$$y_1 = x_0 + x_1$$

$$y_2 = x_0 + x_1 + x_2$$

“Parallel programming is easy as long as you do not care about performance.”

# A Better Parallel Scan Algorithm

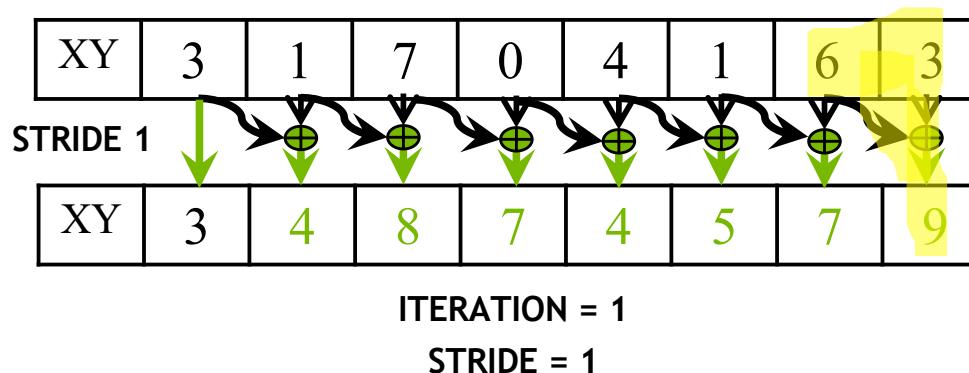
1. Read input from device global memory to shared memory
2. Iterate  $\log(n)$  times; stride from 1 to  $n-1$ : double stride each iteration



- Active threads *stride* to  $n-1$  ( $n$ -stride threads)
- Thread  $j$  adds elements  $j$  and  $j$ -*stride* from shared memory and writes result into element  $j$  in shared memory
- Requires barrier synchronization, once before read and once before write

# A Better Parallel Scan Algorithm

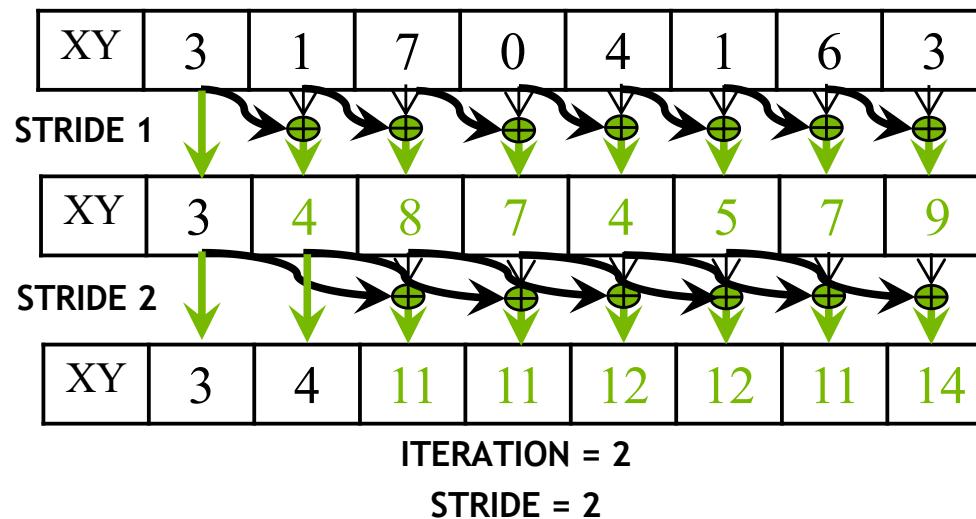
1. Read input from device global memory to shared memory
2. Iterate  $\log(n)$  times; stride from 1 to  $n-1$ : double stride each iteration



- Active threads *stride* to  $n-1$  ( $n$ -stride threads)
- Thread  $j$  adds elements  $j$  and  $j$ -*stride* from shared memory and writes result into element  $j$  in shared memory
- Requires barrier synchronization, once before read and once before write

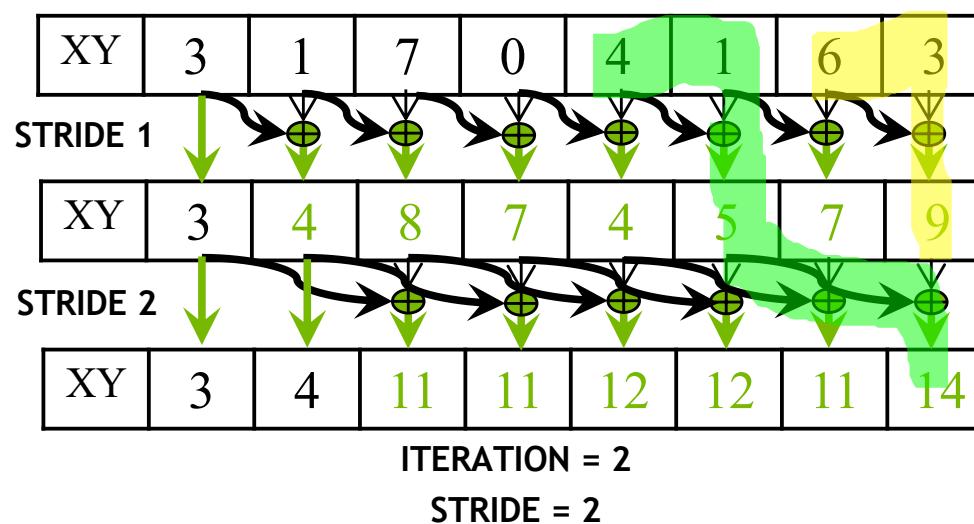
# A Better Parallel Scan Algorithm

1. Read input from device to shared memory
2. Iterate  $\log(n)$  times; stride from 1 to  $n-1$ : double stride each iteration.



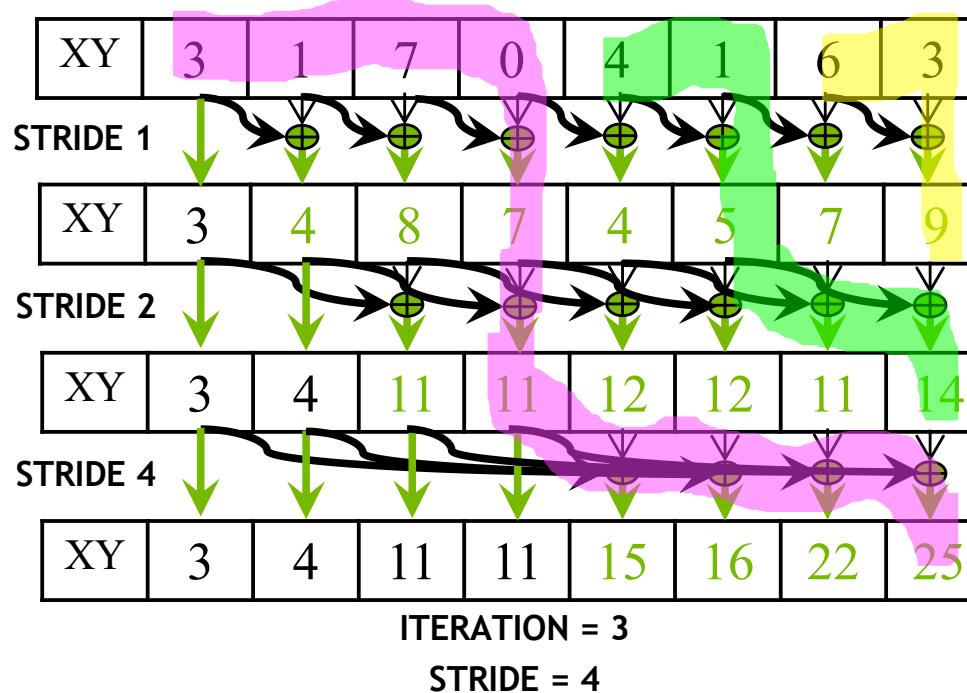
# A Better Parallel Scan Algorithm

1. Read input from device to shared memory
2. Iterate  $\log(n)$  times; stride from 1 to  $n-1$ : double stride each iteration.



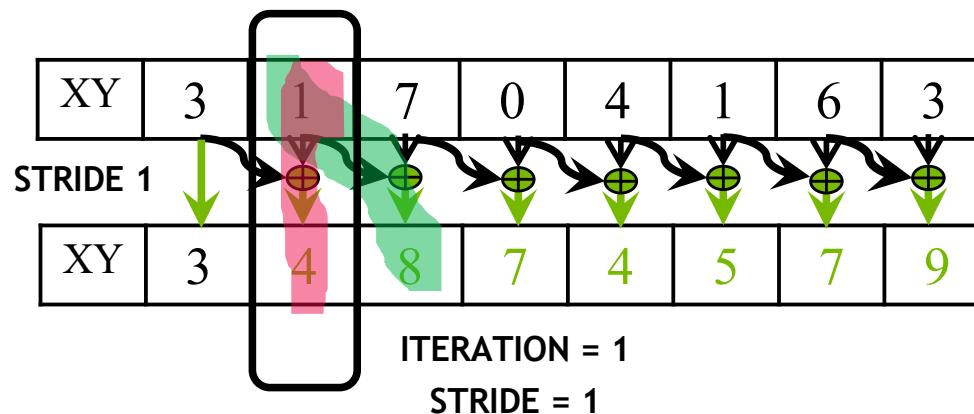
# A Better Parallel Scan Algorithm

1. Read input from device to shared memory
2. Iterate  $\log(n)$  times; stride from 1 to  $n-1$ : double stride each iteration
3. Write output from shared memory to device memory



# Handling Dependencies

- During every iteration, each thread can overwrite the input of another thread
  - Barrier synchronization to ensure all inputs have been properly generated
  - All threads secure input operand that can be overwritten by another thread
  - Barrier synchronization is required to ensure that all threads have secured their inputs
  - All threads perform addition and write output



# A Work-Inefficient Scan Kernel

```
__global__ void work_inefficient_scan_kernel(float *X, float *Y, int InputSize) {
    __shared__ float XY[SECTION_SIZE];
    int i = blockIdx.x * blockDim.x + threadIdx.x;

    if (i < InputSize) XY[threadIdx.x] = X[i]; else XY[threadIdx.x] 0.0f;

    for (unsigned int stride = 1; stride < blockDim.x; stride *= 2) {
        __syncthreads();
        float in1;
        if (threadIdx.x >= stride) in1 = XY[threadIdx.x - stride];
        __syncthreads();
        if (threadIdx.x >= stride) XY[threadIdx.x] += in1;
    }

    __syncthreads();
    If (i < InputSize) Y[threadIdx.x] = in1;
}
```

**SECTION\_SIZE == blockDim.x**  
1 thread per element of the  
data worked on: (n-1) threads  
active in the first iteration of the  
for loop

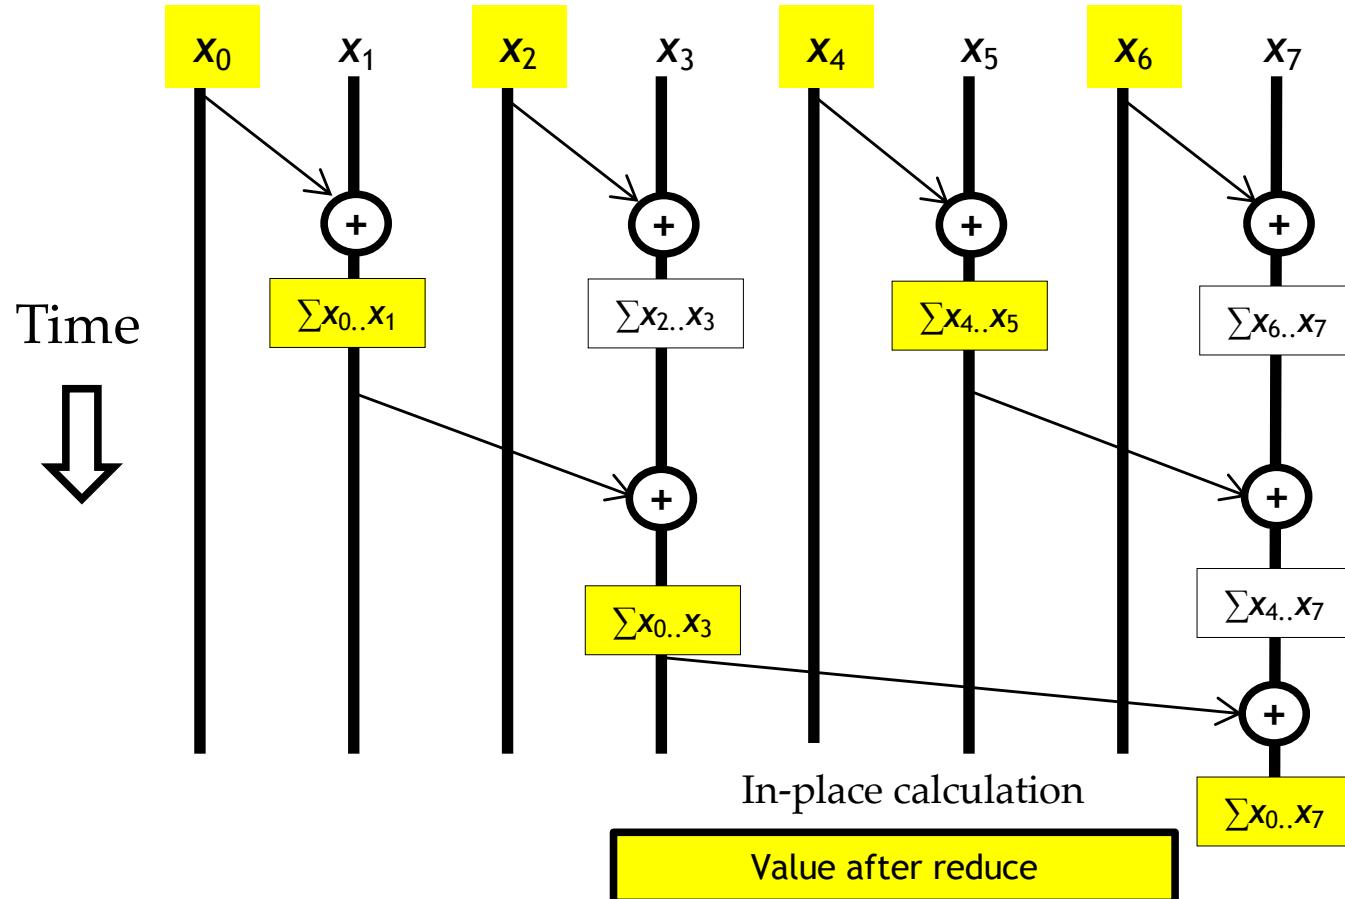
# Work Efficiency Considerations

- This Scan executes  $\log(n)$  parallel iterations
  - The iterations do  $(n-1), (n-2), (n-4), \dots (n - n/2)$  adds each
  - Total adds:  $n * \log(n) - (n-1) \rightarrow O(n * \log(n))$  work
- This scan algorithm is not work efficient
  - Sequential scan algorithm does  $n$  adds
  - A factor of  $\log(n)$  can hurt: 10x for 1024 elements!
- A parallel algorithm can be slower than a sequential one when execution resources are saturated from low work efficiency

# Improving Efficiency

- *Balanced Trees*
  - Form a balanced binary tree on the input data and sweep it to and from the root
  - Tree is not an actual data structure, but a concept to determine what each thread does at each step
- For scan:
  - Traverse down from leaves to the root building partial sums at internal nodes in the tree
    - The root holds the sum of all leaves
  - Traverse back up the tree building the output from the partial sums

# Parallel Scan - Reduction Phase

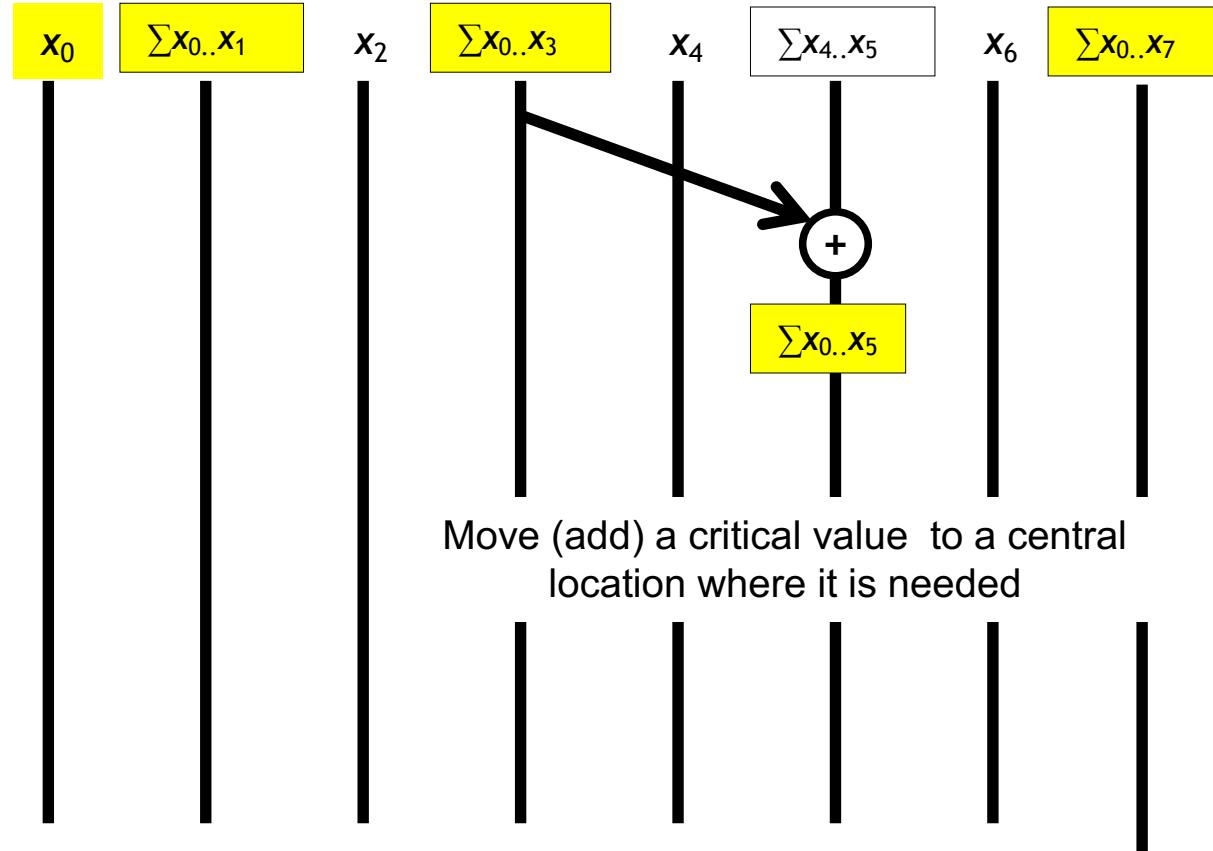


# Reduction Phase Kernel Code

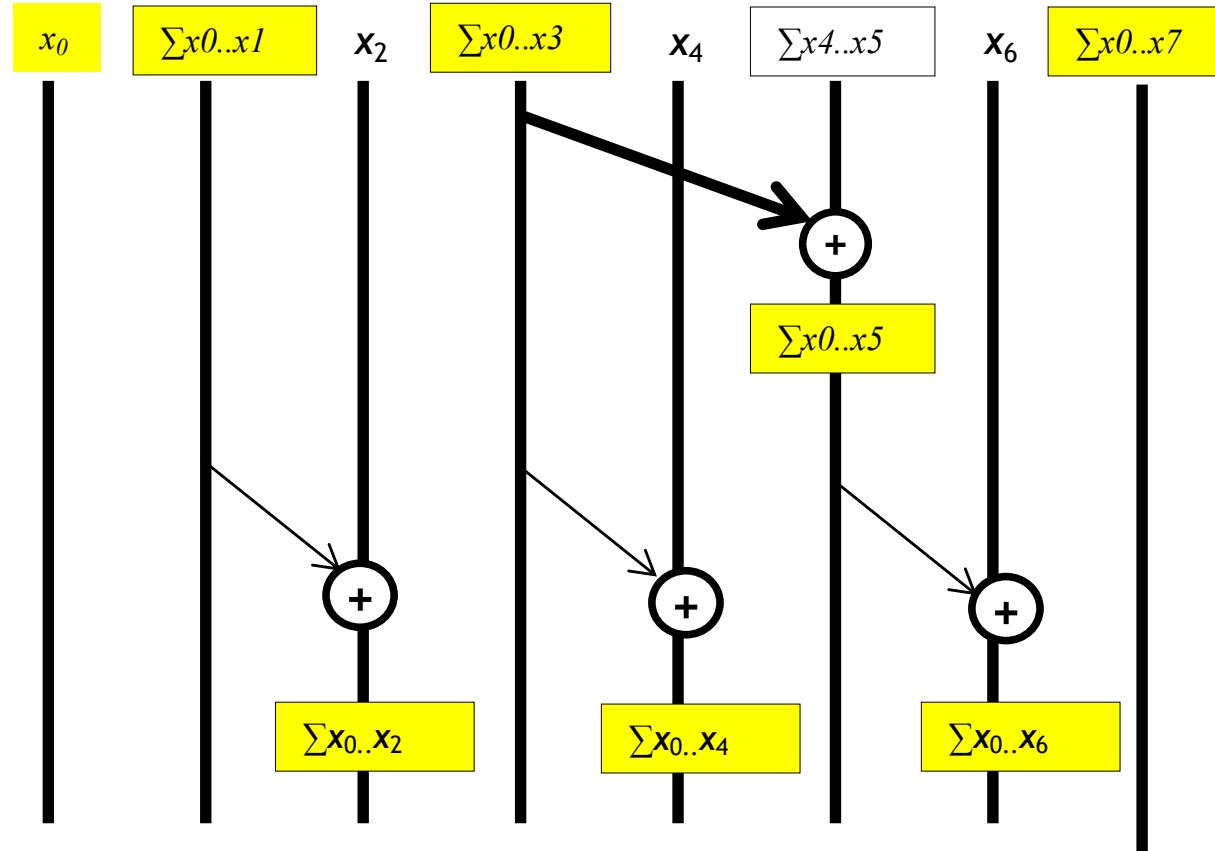
```
// XY[*BLOCKSIZE] is in shared memory SECTION_SIZE = 2*BLOCKSIZE  
  
for (unsigned int stride = 1; stride <= BLOCK_SIZE; stride *= 2) {  
    int index = (threadIdx.x+1)*stride*2 - 1;  
    if(index < 2*BLOCK_SIZE)  
        XY[index] += XY[index-stride];  
    __syncthreads();  
}  
SECTION_SIZE == 2*blockDim.x  
1 thread per two elements of the data worked on:  
threads active in the first iteration of the for loop
```

threadIdx.x+1 = 1, 2, 3, 4....  
stride = 1,  
index = 1, 3, 5, 7, ...

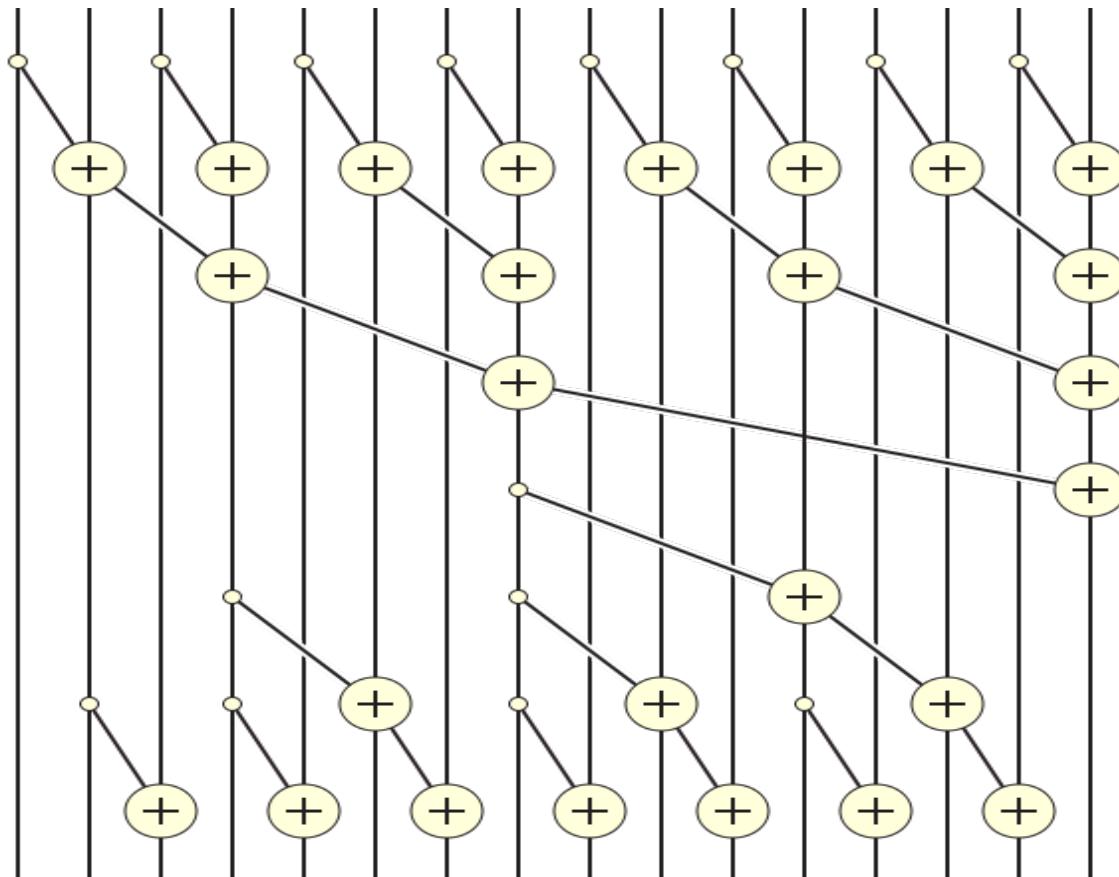
# Parallel Scan - Post Reduction Reverse Phase



# Parallel Scan - Post Reduction Reverse Phase



# Putting it Together



# Post Reduction Reverse Phase Kernel Code

```
for (unsigned int stride = BLOCK_SIZE/2; stride > 0; stride /= 2) {  
    __syncthreads();  
    int index = (threadIdx.x+1)*stride*2 - 1;  
    if(index+stride < 2*BLOCK_SIZE) {  
        XY[index + stride] += XY[index];  
    }  
}  
__syncthreads();  
if (i < InputSize) Y[i] = XY[threadIdx.x];
```

First iteration for 16-element section

BLOCKSIZE=8

stride = BLOCK\_SIZE/2 = 8/2 = 4

For threadIdx.x = 0

index = (0+1) \* (2\*4) - 1 = 8-1 = 7

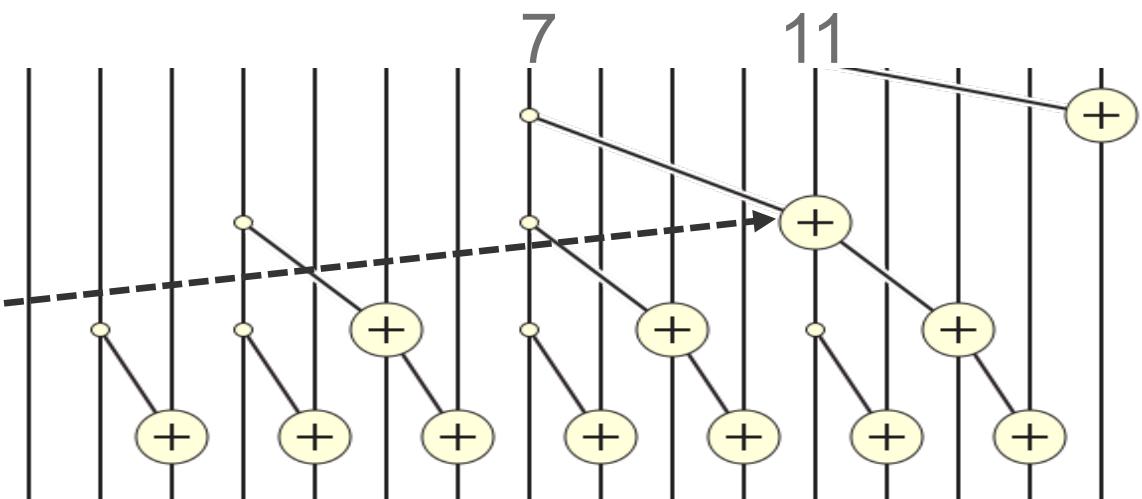
If (7+4< 2\*8) { XY[11] += XY[7] }

For threadIdx.x = 1

index = (1+1) \* (2\*4) - 1 = 15

(15+4 < 2 \*8) NOPE

So only thread 0 gets through on first iteration and updates XY[11] += XY[7]



# Work Analysis of the Work Efficient Kernel

- The work efficient kernel executes  $\log(n)$  parallel iterations in the reduction step
  - The iterations do  $n/2, n/4, \dots, 1$  adds
  - Total adds:  $(n-1) \rightarrow O(n)$  work
- It executes  $\log(n)-1$  parallel iterations in the post-reduction reverse step
  - The iterations do  $2-1, 4-1, \dots, n/2-1$  adds
  - Total adds:  $(n-2) - (\log(n)-1) \rightarrow O(n)$  work
- Together phases perform up to no more than  $2x(n-1)$  adds
- The total number of adds is no more than twice of that done in the efficient sequential algorithm
  - The benefit of parallelism can easily overcome the 2X work when there is sufficient hardware

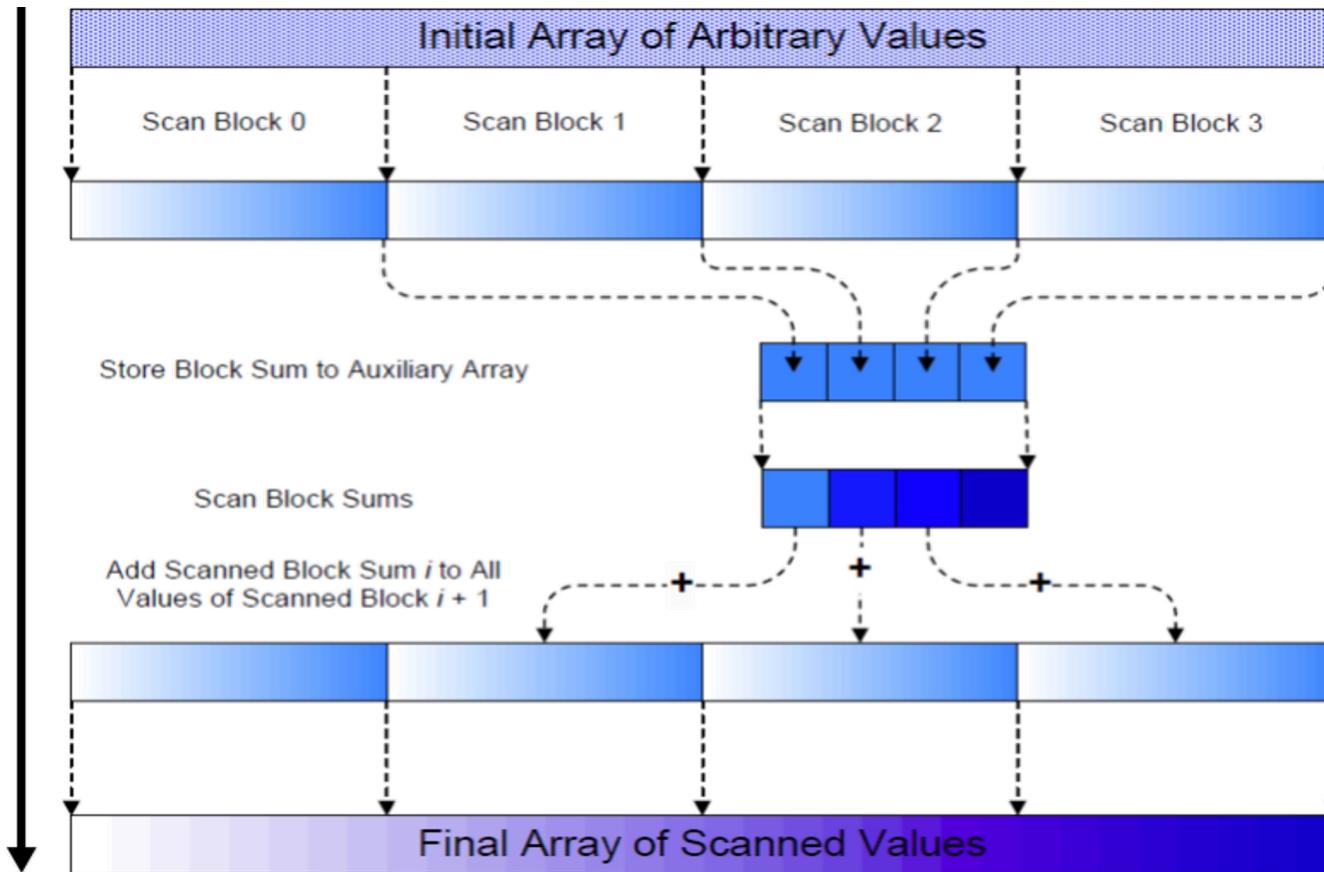
# Some Tradeoffs

- The work efficient scan kernel is normally more desirable
  - Better Energy efficiency
  - Less execution resource requirement
- However, the work inefficient kernel could be better for absolute performance due to its single-phase nature (forward phase only)
  - There is sufficient execution resource

# Handling Large Input Vectors

- Have each section of  $2 * \text{blockDim.x}$  elements assigned to a block
  - Perform parallel scan on each section
- Have each block write the sum of its section into a `Sum[]` array indexed by `blockIdx.x`
- Run the scan kernel on the `Sum[]` array
- Add the scanned `Sum[]` array values to all the elements of corresponding sections

# Overall Flow of Complete Scan



# Exclusive Scan Definition

**Definition:** *The exclusive scan operation takes a binary associative operator  $\oplus$ , and an array of  $n$  elements*

$$[x_0, x_1, \dots, x_{n-1}]$$

*and returns the array*

$$[0, x_0, (x_0 \oplus x_1), \dots, (x_0 \oplus x_1 \oplus \dots \oplus x_{n-2})].$$

**Example:** If  $\oplus$  is addition, then the exclusive scan operation  
on the array [3 1 7 0 4 1 6 3],  
would return [0 3 4 11 11 15 16 22].

# Why Use Exclusive Scan?

- To find the beginning address of allocated buffers
- Inclusive and exclusive scans can be easily derived from each other; it is a matter of convenience

[3 1 7 0 4 1 6 3]

Exclusive [0 3 4 11 11 15 16 22]

Inclusive [3 4 11 11 15 16 22 25]

# A Simple Exclusive Scan Kernel

- Adapt an inclusive, work inefficient scan kernel
- Block 0:
  - Thread 0 loads 0 into XY[0]
  - Other threads load X[threadIdx.x-1] into XY[threadIdx.x]
- All other blocks:
  - All thread load X[blockIdx.x\*blockDim.x+threadIdx.x-1] into XY[threadIdx.x]
- Similar adaption for work efficient scan kernel but ensure that each thread loads two elements
  - Only one zero should be loaded
  - All elements should be shifted to the right by only one position

Read the Harris article (Parallel Prefix Sum with CUDA) for a more intellectually interesting approach to exclusive scan kernel implementation.



## GPU Teaching Kit



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the [Creative Commons Attribution-NonCommercial 4.0 International License](#).